

# *BlkKin*: An End-to-end Tracing Tool for Software Defined Storage Systems

Marios-Evangelos Kogias<sup>1</sup>, Filippos Giannakos<sup>1</sup>, Vangelis Koukis<sup>1</sup>, Constantinos Venetsanopoulos<sup>2</sup>, and Nectarios Koziris<sup>1</sup>

<sup>1</sup>National Technical University of Athens

<sup>2</sup> okeanos grnet

## Abstract

Distributed storage systems require special treatment concerning their monitoring and tracing. The need for production-scale, live tracing arises since under these circumstances, interesting and problematic behaviors are most likely to become obvious. However, this need in conjunction with low-latency and high-throughput demands set further challenges, apart from the initial ones concerning distributed tracing like clock synchronization and trace correlation. BlkKin is a tracing infrastructure used for software-defined storage systems' live tracing. It inserts minimal overhead to the instrumented system and it provides a Web UI for data visualization. BlkKin gives us an end-to-end system representation and metrics which enables us to extract information on how does the system performs, either per module or as a whole, per request, under different workloads.

**Keywords:** tracing, monitoring, low-latency, LTTng, RADOS

## 1 Introduction

The more complex the distributed systems are becoming today, the greater grows the need for better debugging and monitoring tools. Debugging or monitoring a distributed system is a difficult and demanding job. Dealing with multiple hosts makes the system's behavior unpredictable and bound to a specific context that could change from execution to execution. Thus, finding failures or bottlenecks cannot be achieved through traditional monitoring

or debugging tools like syslog or gdb.

This problem can be tackled through tracing. Tracing captures the state of the system along with other needed information that could provide further insight on the exact context in which a specific event happened. The information can be correlated to explore the system's behavior under different working states. Unlike traditional monitoring, to obtain this kind of information, instrumentation points need to be placed in the system's source code indicating the events to be traced and the values to be measured. Through this instrumentation we want to trace **causal relationships** between distinct software layers, **latencies** and **bottlenecks**, as well as unknown **execution paths**.

Distributed tracing sets the following challenges. The system should be traced under real working conditions, since most of the failures are more likely to appear at full load. So, ideally we want to have a fully functional system working at production scale that is getting traced at the same time. Consequently, the tracing mechanism should not affect the system's performance significantly. Also, we have to deal with multiple hosts and events happening almost simultaneously. We need a mechanism that qualitatively correlates distributed event information, namely a way to group correlated, tracing information. Finally, we need to handle with the absence of a global clock.

Some of the most widespread distributed systems are the distributed storage systems. They can be used either as huge data warehouses for content like images or videos, or as storage backends for virtual volume provisioning employed by IaaS providers.

In any case, their performance analysis and tuning is of vital importance for better resource usage and faster, responsive services.

In this paper, we present the design and implementation of *BlkKin*. BlkKin is a system that enables us to debug and monitor through tracing a distributed storage system in real time, with very low overhead and visualize the aggregated information. To build BlkKin we combine various open-source technologies. We use LTTng (Linux Trace Toolkit - next generation)<sup>1</sup> as tracing backend per host and we extend it to communicate in real time with Zipkin<sup>2</sup>, a Twitter’s, open-source implementation of Google’s Dapper[10]. We created a C/C++ library to instrument applications according the semantics described in [10] and using this library we extended Archipelago[6] and RADOS[12] to support cross-layer tracing using BlkKin. Basically, this means that an IO request is associated with a specific id which is propagated within the infrastructure until the request is finally served. So, we end up with a visualization of the request’s route from the time it is initiated within a virtual machine until it is served by RADOS, accompanied with meaningful information like request type or the target object’s name.

Other approaches, like Ganglia[8], were more monitoring oriented while others like Stardust[11] or [3], did not provide live tracing.

## 2 Background

In this section we describe why we used Zipkin and LTTng as BlkKin’s building blocks.

### 2.1 Tracing logic

According to [10] there are two dominant schools in associating all record entries with a given initiator. *Black box* monitoring scheme assumes there is no additional information other than the message record. So statistical regression techniques should be used to infer any existing association. *Annotation-based* schemes, though, rely on applications or middleware to explicitly tag every record with a global identifier that links these message records back to the originating request.

<sup>1</sup><http://lttng.org/>

<sup>2</sup><http://twitter.github.io/zipkin/>

For our needed, overall review of the system per specific request, we had to implement an annotation-based monitoring scheme. In addition, this scheme needs to depict the causal relationships between the different layers, taking into account that the system is distributed, thus providing information about the hosts. Finally, it should enable us to collect any other information considered important and not restrict us to time data.

Google proposed a complete annotation-based scheme in [10]. This scheme meets the aforementioned demands and is used in BlkKin. [10] describes the following concepts for tracing:

**annotation** The actual information being logged. There are two kinds of annotations. Either *timestamp*, where the specific timestamp of an event is being logged or *key-value*, where a specific key-value pair is being logged.

**span** The basic unit of the process tree. Each specific processing phase can be depicted as a different span. Each span should have a specific name and a distinct span id.

**trace** Every span is associated with a specific trace. A different trace id is used to group data so that all spans associated with the same initial request share the common trace id. For our case, information concerning a specific IO request share the same trace id and each distinct IO request initiates a new trace id.

**parent span** In order to depict the causal relationships between different spans in a single trace, parent span id is used. Spans without a parent span ids are known as *root spans*.

So, by creating tracing data according to these semantics we can have an end-to-end sense of our system’s performance, behavior and internal latencies that may vary depending on the nature and size of each request which can also be captured using key-value annotations.

### 2.2 Logic implementation

Based on these primitives, Twitter created Zipkin, a distributed tracing system used to gather timing data for all the disparate services running on their premises. Zipkin consists of a data collector,

a database service, with support for a variety of SQL and NoSQL databases, and a Web interface to visualize the aggregated data.

Zipkin uses Scribe[1] to transport traces from the different services to the central collector. Scribe is a logging server created by Facebook, aiming to be scalable and reliable. Scribe servers are arranged in a directed graph, with each server knowing only about the next server in the graph. This network topology allows for adding extra layers of fan-in, as a system grows and batching messages before sending them between datacenters as well as providing reliability in case of intermittent connectivity or node failure. Scribe makes use of the Thrift protocol[2] for data transfer.

Zipkin seemed to fit our demands concerning data collection since it is designed to scale. Also, apart from the Web UI, we could execute ad-hoc queries on the database for more elaborate information extraction, without the need for batch processing Map-Reduce jobs. However, Zipkin did not provide any C/C++ library to instrument low-overhead applications. The offered libraries were in scripting languages like Ruby, or in Java and Scala. We created a C/C++ library that encapsulates the Dapper semantics and can be used within C/C++ projects to create trace information in accordance with the Dapper logic. This library is designed to be backend-independent, which means that one can implement his own log aggregation backend. However, we offered a specific backend implementation according to our initial prerequisites concerning overhead, based on LTTng that is being thoroughly examined in the next chapter.

## 2.3 Tracing Backend

BlkKin’s design was driven by a strict low-overhead and production-wise operation prerequisite. So the first important decision to make concerned the system that would implement the system’s backend, namely the system that would run on every cluster node and be responsible for aggregating tracing data from the instrumented applications. We chose to use LTTng in our system’s backend.

LTTng is a toolchain that allows integrated kernel and user-space tracing from a single user interface. It was initially designed and implemented to reproduce, under tracing, problems occurring in normal conditions. It uses a linearly scalable and

wait-free RCU (Read-Copy Update) synchronization mechanism and provides zero-copy data extraction. These mechanisms were implemented in kernel and then ported to user-space as well. In addition, LTTng supports a variety of Linux distributions and since version 2.x kernel tracer modules are built against a vanilla or distribution kernel, without need for additional patches. The nanosecond-scale added overhead per event traced mentioned in [5], convinced us to use LTTng and have the same generic toolkit for both user and kernel tracing, which can be easily installed on almost every system needed to instrument.

Concerning user-space tracing, LTTng supports static tracepoint instrumentation. This means to manually insert tracepoints in the application source code and rebuild the application. After rebuilt, these tracepoints will get triggered during execution and produce the described traces, without breakpoints or system-calls. As far as kernel-space is concerned, LTTng is supported by kprobes[9] and kernel markers[4]. Thus LTTng does not significantly affect the system’s performance. Although the whole process of recompiling the instrumented application may seem counterproductive, static instrumentation abilities are limitless. Based on the knowledge and understanding of the application, one can instrument and trace every part that might be problematic or causing longer latencies, as well as extracting all the information needed to fully understand under which context each request was served. Consequently, since we chose to implement the Dapper tracing semantics, static tracing was the only way to do that.

Finally, live tracing is a prerequisite. Although live monitoring may seem natural for systems like ELK<sup>3</sup>, live tracing in combination with low-overhead was not possible. Older tracing approaches required separation between the tracing and operation phases. This happened either because tracing added a lot of extra overhead to the instrumented system so it could not continue working production-wise, or because tracing information was not in a state that could be processed before the end of the tracing session. For LTTng especially, tracing data were in a binary form that could be decoded only after the end of the tracing session. However, since version 2.4 of LTTng there is also

---

<sup>3</sup><http://www.elasticsearch.org/overview/>

live tracing support. Using a daemon called *relayd*, tracing data can be streamed over TCP forming batch TCP packets and processed while the tracing session is still in progress.

### 3 Design

As mentioned, BlkKin was designed respecting four prerequisites: **low-overhead tracing**, **live tracing**, **Dapper semantics** and **Web UI**.

#### 3.1 BlkKin architecture

So far we mentioned that we need a tracing agent that run on every host and captures data produced by the instrumented applications. This agent is LTTng. Then, data have to be sent in real time to Zipkin. However, LTTng live tracing supports only CTF<sup>4</sup> to text transformations. Consequently, to process and visualize tracing data in real time in Zipkin, we implemented a live-tracing plugin that transforms in real time CTF data to Scribe messages and sends them to a Scribe server. Our plugin, base on Babeltrace<sup>5</sup>, which is a CTF converter and trace viewer, reads and decodes the CTF-encoded information, creates Thrift encoded messages recognizable by the Zipkin collector and sends these messages to the Scribe server. The plugin is under evaluation for contribution to LTTng.

So we BlkKin's architecture is presented in Figure 1. The instrumented application produces tracing information according to the Dapper semantics using our BlkKin library, through instrumentation points in its source code. This information is aggregated by LTTng and sent to the relayd. After that, our Babeltrace plugin gets the tracing data from the relayd, processes them as mentioned and finally sends them to the Scribe server.

##### 3.1.1 BlkKin Deployment

Each arrow show in Figure 1, in BlkKin's internal architecture represents a TCP communication. Also, the Scribe server, which the Babeltrace plugin finally sends the tracing information to, can be either the central Zipkin collector or a local Scribe daemon that will eventually send the data to the

central collector, thus exploiting the asynchronous, buffered communication provided by Scribe. So we can have a lot of versatility concerning deployment.

In a cluster deployment with many nodes, we end up having the architecture presented in Figure 1. In this deployment, there is a whole BlkKin stack running on every cluster node. All the communication takes place over localhost and the tracing data will end up being handled by a local Scribe daemon. This daemon will normally send the data to the central Zipkin collector in batch messages, but also store them locally in case of a connection loss or if the central Zipkin collector is busy and try to send them later. So through batch messaging we reduce the Zipkin collector congestion and we can make sure that no data will be lost, because any data loss would mess up the tracing semantics and end up with an inconsistent UI and database state.

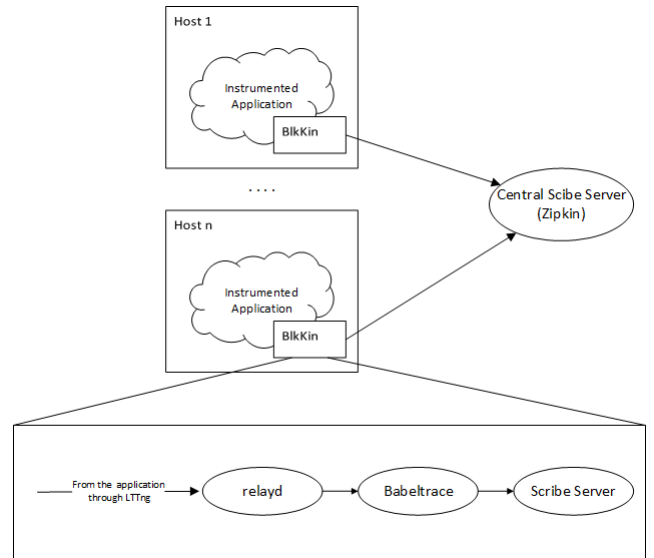


Figure 1: BlkKin deployment and communication

#### 3.2 Clock Synchronization

A matter that needs special treatment when it comes to distributed tracing, is clock synchronization. In a BlkKin cluster deployment, ideally we would like to have no time difference between the clocks of individual cluster nodes, since we are interested in measuring latencies between nodes and processing durations in the scale of microseconds. So, a possible time skew in the scale of milliseconds

<sup>4</sup><http://www.efficios.com/ctf>

<sup>5</sup><https://ltnng.org/babeltrace>

could result in having a response virtually happening before its triggering request. The most common solution for clock synchronization is NTP. Although according to [3] NTP’s accuracy was not acceptable for tracing, so arithmetic methods had to be deployed to solve the time skew problem, the NTP v4 published in 2010, improved NTP’s potential accuracy to the tens of microseconds with modern workstations and fast LANs, through fundamental improvements in the mitigation and discipline algorithms. So NTP’s accuracy fit our needs for clock synchronization.

## 4 Implementation

Archipelago is a software-defined storage system, used in Synnefo[7] and uses RADOS as a storage backend. Archipelago is structured in a layered architecture. It offers a QEMU driver for provisioning virtual volumes in virtual machines. So using our BlkKin library we instrumented the QEMU driver, Archipelago and RADOS source code.

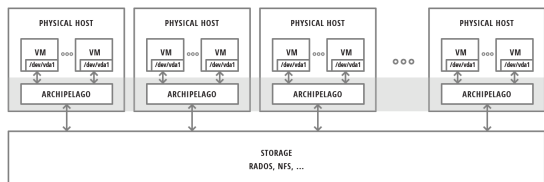


Figure 2: Archipelago overview

Imitating other Zipkin libraries that used HTTP headers to transfer the trace, span and parent span ids, our library used a struct containing this information. This struct is initiated in QEMU and propagated through the Archipelago layers, as a separate field in Archipelago requests, and finally through librados to RADOS. By extending the RADOS internal classes, we managed to reach this information until the request is served by the filestore. Overall, we added 113 instrumentation points, that log the separate phases the IO request passes through and all the needed information concerning the specific request, like the object name, the request type (read/write) and its size.

## 5 Evaluation - Conclusion

We ran benchmarks to evaluate the added overhead that BlkKin imposes on Archipelago and RADOS. For example imitating an IOZone workload we measured the throughput of 1GB of 64KB sequential writes without instrumentation points and with instrumentation points with enabled and disabled tracing, since LTTng can stop the tracing session without affecting the application and start it again if needed. The results can be seen in the table.

Mode	Throughput(MB/s)	Latency/Op (ms)
not instrumented	12.7	5.16
instrumented stopped	11.8	5.56
instrumented	11	5.94

So, the layered multi-node architecture shown in Figure 2, after being traced with BlkKin, ends up as shown in Figure 3 in the Zipkin UI. Each bar represents a distinct processing phase.

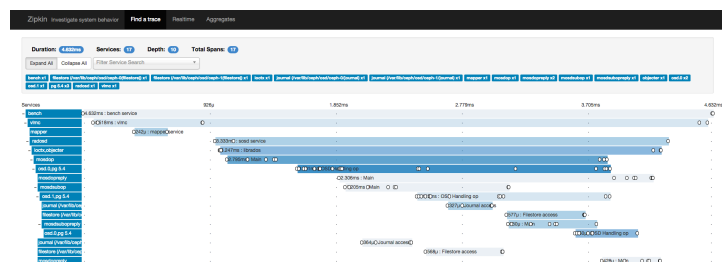


Figure 3: Archipelago overview

Apart from the UI that gave us an overall impression about the systems’ performance and bottlenecks, using simple SQL-queries we were able to calculate average metrics like each layer’s duration, the communication latency between Archipelago and RADOS or the time that a request stayed in RADOS’s dispatch queue before being served.

To sum up, we created a working prototype for a low-overhead, live-tracing infrastructure based on open-source technologies. Using BlkKin we can extract information not only per system’s component but also, per request. Exposing the system’s internals under any possible workload can lead to more reliable distributed systems.

## References

- [1] Data warehousing and analytics infrastructure at Facebook. *Special Interest Group on Management of Data (SIGMOD)*.
- [2] Aditya Agarwal, Mark Slee, and Marc Kwiatkowski. Thrift: Scalable cross-language services implementation. Technical report, Facebook, 4 2007.
- [3] Eric Anderson, Christopher Hoover, Xiaozhou Li, and Joseph Tucek. Efficient tracing and performance analysis for large distributed systems. In *MASCOTS*, pages 1–10. IEEE, 2009.
- [4] Jonathan Corbet. Kernel Markers. In *Linux Weekly News*. August 2007.
- [5] M. Desnoyers. *Low-impact Operating System Tracing*. PhD thesis, Ecole Polytechnique de Montreal, 2009.
- [6] Filippos Giannakos, Vangelis Koukis, Constantinos Venetsanopoulos, and Nectarios Koziris. okeanos large-scale cloud service using cepg. *login*, 39(3):6–9, June 2014.
- [7] Vangelis Koukis, Constantinos Venetsanopoulos, and Nectarios Koziris. Synnefo: A complete cloud stack over ganeti. *login*, 38(5):6–10, October 2013.
- [8] Matthew L. Massie, Brent N. Chun, and David E. Culler. The ganglia distributed monitoring system: Design, implementation and experience. *Parallel Computing*, 30:2004, 2003.
- [9] Ananth Mavinakayanahalli, Prasanna Panchamukhi, Jim Keniston, Anil Keshavamurthy, and Masami Hiram. Probing the guts of kprobes. In *Proceedings of the 2006 Linux Symposium*, Ottawa, Canada, July 2006.
- [10] Benjamin H. Sigelman, Luiz Andr Barroso, Mike Burrows, Pat Stephenson, Manoj Plakal, Donald Beaver, Saul Jaspan, and Chandan Shanbhag. Dapper, a large-scale distributed systems tracing infrastructure. Technical report, Google, Inc., 2010.
- [11] Eno Thereska, Brandon Salmon, John D. Strunk, Matthew Wachs, Michael Abd-El-Malek, Julio Lpez, and Gregory R. Ganger. Stardust: tracking activity in a distributed storage system. In Raymond A. Marie, Peter B. Key, and Evgenia Smirni, editors, *SIGMETRICS/Performance*, pages 3–14. ACM, 2006.
- [12] Sage A. Weil, Andrew W. Leung, Scott A. Brandt, and Carlos Maltzahn. Rados: a scalable, reliable storage service for petabyte-scale storage clusters. In *Proceedings of the 2nd international workshop on Petascale data storage: held in conjunction with Supercomputing '07, PDSW '07*, pages 35–44, New York, NY, USA, 2007. ACM.